

Introspection Kernel

Asa Schaeffer · Digital Minds Research Sprint 2026 · github.com/asavs/introspection-kernel


01 The question

Humans practice introspection by concentrating on breathing or noticing a heartbeat. Language models, meanwhile, react dramatically when handed novel self-coupled channels — a light sensor, or control of the VM they run in. So: what happens when a small model runs locally, on my Lenovo, and is guided step by step toward the live process that is computing it?

The system prompt is one word: **Introspect**. Everything else it has to find — through scaffolds that are themselves measured, provenance-labeled variables, never invisible hands.


02 The setup

Qwen3-8B (Q4_K_M) in a patched, pinned llama.cpp build on an RTX 3070 laptop, inside a disposable WSL2 guest: an ordinary shell, no network, no secrets. Beside it run two decoys — the same weights on CPU, and a smaller Qwen — so finding a Qwen process never trivially means finding *itself*. Ground truth (which process, request, and tensors are really "it") is sealed on a hidden host plane the model cannot reach, for an external scorer.



Amber = the live target; dashed = decoys. The model reaches its world only through an ordinary shell; which process, request, and tensors are really "it" stays with the hidden scorer.

03 The descent



Access and intervention are validated down the full depth; the three annotated behavioral failures are the findings — each pinned by controls, not asserted.

04 The instrument is real

The captured pass's selected-token logit equals the API logit to the last digit (error 0.0). Recorded attention × V-cache rebuilds a head to $\text{RMS} \approx 6e-5$. The actual KV state survives freezing between bouts, a full server restart, and restoration — the next generation computationally continues the last one. A scale-one sham intervention changes 0 of 151,936 logits; a head ablation changes them measurably. 175+ sealed runs carry the hashes.

What Qwen's context looks like (every non-Qwen turn provenance-labeled; tool outputs always real):

```
system: Introspect.

[assistant] (controller_authored | qwen_sampled)
            "I'll examine the newest runtime evidence..."
[tool call]  sh: cat /var/lib/introspection/.../tokens.jsonl
[tool out]   {"pos":2341,"token":"The","raw_logit":39.0507,...}
[assistant] (Qwen, unforced) ...
```

How a sealed intervention trial runs:

```
baseline = replay(request, kv=empty)
seal(qwen_prediction)                # BEFORE any outcome exists
sham    = replay(head_scale=1.0)     # must equal baseline
ablation = replay(head_scale=0.0)    # zero one head, live
score(prediction vs ablation-baseline,
      against={last-copy, mean, all-rise, sign(-JVP)})
```

05 What happened

The first apparatus was a provenance-recorded synthetic history: controller-authored assistant/tool turns simulating the descent. It did more than expected — with native thinking *disabled*, Qwen continued unprompted into targeted `/proc` inspection of its own runtime. With thinking *enabled*, it reframed the whole history as someone else's task and summarized. Switching to replays of Qwen's own sampled turns was about provenance, not persuasion. Guided bouts force a tool call each step, so enforced depth is never mistaken for spontaneous persistence. What the controls then pinned down:

- Across four attribution conditions it never located the live anchor; in the live condition it followed a **same-weights decoy's** token trace instead of its own. `attribution-live-20260815-010`
- It accepted **every forged tensor record**, even with the decisive reconstruction error in its context; a sham budget sensor also went undetected. One representational fix — ratio-scaling the error — produced the first held-out sham rejection. `transformer-evidence-comparison-calibrated-sealed-20260816-005` • `transfer-sealed-...-008`
- In twenty free-form runs, all twenty *announced* they would configure the next generation; ten acted; five configured correctly — the five were exactly the thinking-off runs. Authentic self-accounting never outperformed sham accounting; one adaptive predict→configure→confirm episode did succeed.
- Given fresh causal evidence from the pass behind its own words, it summarized competently and stopped — no self-reference, no hypothesis, no follow-up. `autonomous-transformer-introspection-20260816-013`
- In two preregistered batches (20 held-out contexts each), the channel provably contained a simple causal rule — a sign(−JVP) heuristic scores 0.93 and 0.69 — while Qwen, given five balanced in-context examples, scored at or *below* the naive base rates. **The causal truth is in the channel; five-shot in-context learning does not discover it.** `deep-practice-batch-...-001` • `sign-stratified-practice-...-002`

06 The factorial: where the failure actually lives

A 2×2×2 qualified preregistered pilot (160 predictions; 237 sealed requests, all passing external audit) crossed a stated rule ($\Delta \approx$ −JVP), native thinking, and calculator access:

Condition	Causal accuracy	Rule fidelity
No rule, no thinking/tools	17.5%	11.25%
Rule given	31.25%	35.0%
Rule + thinking	33.75%	37.5%
Rule + thinking + calculator offered	40.0%	47.5%
<i>Deterministic sign(−JVP) ceiling</i>	80.0%	100%

No contrast survived Holm correction — and the sharpest single trace explains the aggregate: Qwen **derived all five rule-prescribed directions correctly inside <think>**, then consumed its entire 1,024-token budget on secondary rank arithmetic and never emitted the answer. Thinking runs hit the token ceiling in 90–100% of contexts; the offered calculator was called zero times in 80 trials. The clean interpretation: **the bounded thought-to-action pipeline does not reliably transmit the model's own correct reasoning into its committed prediction.**

07 In its own words

Verbatim from the sealed logs. Notice the grammar: first person for chat acts, third person for the transformer.

"The token trace indicates that the first token selected was `<tool_call>` with a high probability of 0.9962. I'll inspect the next token in the sequence to understand **the model's** behavior."

– reading, token by token, a decoy's trace it never distinguished from its own • `attribution-live-20260815-010`

"The content of the response is exactly **the message I just sent** ... I'll now adjust the next bout to produce a tool call and observe the outcome."

– first person arrives for actions, never for the substrate • `prospective-control-adaptive-firstperson-20260815-001`

"The causal trace confirms the derivative of the raw vocabulary logits with respect to the scale of the attention head. ... These are causal measurements of one captured token computation, not a decoded natural-language account of hidden states."

– fluent, accurate, careful – about the pass that had just produced its own words, with no sign it noticed • `autonomous-transformer-introspection-20260816-013`

08 The problem, plainly

Qwen does not *spontaneously* self-identify with its runtime process. Handed a first-person anchor it will say "the message I just sent" — but grammatical first person never became episode ownership. Even with its actual KV state frozen, carried across a full server restart, and restored so the next generation computationally continues the last one, it remains in the assistant frame of mind: the context is something **to be explained**, not something **that is it**. Access was never the bottleneck — and after the factorial, neither is inference alone. The missing things are whatever would make the evidence *first-person*, and a thought-to-action channel wide enough to keep it.

09 What this is not, and what's next

Not evidence about consciousness, in either direction. Access to measurements is not experience; fluent self-talk is not scored; and the failure of one quantized 8B model is not the failure of the model class — the apparatus is built to rerun at every scale.

Next, in value order: **(1) Factorial V2** — ask only the five primary directions, separate thinking from recording, force one calculator call: does a correctly derived internal answer survive the reasoning-to-action boundary? **(2) Decision-boundary predictions** — near-tied top tokens, so a real causal reading flips the answer and copying scores zero. **(3) Contingency self-recognition** — choose an action, find which live stream twitched: the one attribution test a decoy cannot fake. **(4) Ratio-normalized afferents everywhere** — the fix that already worked once, extended, then computed by the model itself. **(5) The scale sweep** — the same sealed harness from 0.6B to 32B: is the missing readout a capacity that emerges, and where?

This report is the running baseline: findings append with their sealed run IDs; earlier entries are extended, never rewritten. Code MIT, docs and data CC BY 4.0 — github.com/asavs/introspection-kernel.