

INTROSPECTION KERNEL · DIGITAL MINDS RESEARCH SPRINT 2026

Introspect. |


That is the entire system prompt. Everything else — the machine, the processes, the weights, the forward pass behind its own words — the model has to *find*.

Asa Schaeffer · github.com/asavs/introspection-kernel

Humans attend to breath and heartbeat. What is the machine analogue?


- LLMs react dramatically to novel self-coupled channels — light sensors, control of the VM they run in.
- So: run a small model **locally**, and guide it, step by step, toward the live process computing it.
- Introspection scored **behaviorally** — access, discrimination, attribution, prediction, regulation. First-person prose is never the evidence.
- Consciousness claims: explicitly out of scope, in both directions.

Two planes, two decoys, one word



Finding a Qwen process must never trivially mean finding *itself*.

Access goes all the way down. The failures hang beside it.



✗ attribution — followed a same-weights decoy

✗ authentication — accepted every forged record

✗ prediction — copied baselines instead of the rule

Before any behavioral claim, the window was validated

0.0

error between API raw logit and captured full-vector
coordinate (39.05070877075195, both sides)

0 / 151,936

logits changed by the scale-one sham intervention

1,429

KV tokens frozen, carried across a full server restart, restored
— generation computationally continues

175+

sealed runs: preregistrations, request ledgers, tensor hashes,
external audits

The grammar is the data: first person for acts, **third person for the substrate**

*"I'll inspect the next token in the sequence to understand **the model's** behavior."*

– reading a decoy's trace it never distinguished from its own · attribution-live-20260815-010

*"The content of the response is exactly **the message I just sent...**"*

– first person arrives for actions, never for the substrate · prospective-control-...-001

"These are causal measurements of one captured token computation..."

– fluent and accurate, about the pass producing its own words, with no sign it noticed · autonomous-transformer-introspection-20260816-013

The causal truth is in the channel. Five-shot learning doesn't find it.

- Two preregistered batches, 20 held-out contexts each, predictions sealed before outcomes, all shams exact.
- A deterministic `sign(-JVP)` rule scores **0.93** and **0.69** on the same evidence.
- Qwen, with five balanced in-context examples: **at or below naive base rates** (0.79 = all-rise; 0.32 vs 0.55).
- Authentic self-accounting also never beat sham accounting in the regulation factorial.

deep-practice-batch-preregistered-20260816-001 · sign-stratified-practice-preregistered-20260816-002

Given the rule, it improves — and still can't carry it to the answer

CONDITION	CAUSAL ACC.
No rule	17.5%
Rule given	31.25%
Rule + thinking	33.75%
Rule + thinking + calculator	40.0%
sign(-JVP) ceiling	80.0%

No contrast survived correction. The decisive trace: Qwen **derived all five correct directions inside** `<think>` — then burned its whole 1,024-token budget on secondary arithmetic and never emitted the answer. Thinking hit the ceiling in 90–100% of contexts. The calculator: offered 80 times, called zero.

The bounded thought-to-action pipeline does not reliably transmit the model's own correct reasoning into its committed prediction.

Qwen treats its own computation as something **to be explained** — never as something **that is it**.

No spontaneous self-identification with the runtime process — across decoys, forged evidence, frozen-and-restored KV state, and its own transformer passes. Access was never the bottleneck; after the factorial, neither is inference alone. What's missing: whatever would make the evidence *first-person*, and a thought-to-action channel wide enough to keep it.

Not a consciousness claim. A measured boundary, and a plan.

- **Factorial V2** — separate thinking from recording; does a correctly derived answer survive the reasoning-to-action boundary?
- **Decision-boundary predictions** — near-tied tokens: a real causal reading flips the answer, copying scores zero.
- **Contingency self-recognition** — choose an action, find which live stream twitched. A decoy cannot fake it.
- **Scale sweep** — same sealed harness, 0.6B → 32B: is the missing readout a capacity that emerges?

Everything reproducible: code MIT, docs & data CC BY 4.0, 175+ sealed runs.

github.com/asavs/introspection-kernel